

IST 707 – Data Analytics
Black Friday Sales
Group 6

Kshitij Sankesara

Vijet Muley

Yunxi Dong

Introduction:

A retail store wants to know which products to stock before the holiday season to make sure that the store doesn't run out of products. For this, the store wants to analyze all the previous year's transactions. They want to make sure that they don't over stock or under stock any particular category of product.

To help the store owners we analyze the complete Black Friday dataset which is a sample of the transactions made in the retail store. We get to know the customer purchase behavior against different products.

The report will showcase how we have gone through with our analysis, starting from data cleaning, preprocessing the data, descriptive statistics, visualizations, generating models for the analysis, comparison of our models and conclusion based on all the whole analysis.

Data Description:

The Black Friday Dataset was obtained from Kaggle which consists of 538k rows and 12 attributes. The attributes are User_ID, Product_ID, Gender, Age, Occupation, City_Category, Stay_in_Current_City_Years, Marital_Status, Product_Category_1, Product_Category_2, Product_Category_3 and Purchase.

Variable	Definition
User_ID	User ID
Product_ID	Product ID
Gender	Sex of User
Age	Age in bins
Occupation	Occupation (Masked)
City_Category	Category of the City (A,B,C)
Stay_In_Current_City_Years	Number of years stay in current city
Marital_Status	Marital Status
Product_Category_1	Product Category (Masked)
Product_Category_2	Product may belongs to other category also (Masked)
Product_Category_3	Product may belongs to other category also (Masked)
Purchase	Purchase Amount (Target Variable)

Attributes and their Definitions

Data Preprocessing:

First, we found out the missing values in our dataset. There were missing values in two columns 'Product_Category_2' and 'Product_Category_3'. We tried different methods for the missing values but later we decided to replace all the missing values with 0. We replaced it with 0 as it made more sense for our later analysis.

It took a while for us to understand as to why this may be the case. Rest of the data has no missing value whatsoever. However, only these columns seem to display missing values. The answer to this was actually in the data description. The product category columns are present in order to cover the versatility of a product. Meaning, the product category bins in 'Product_Category_1' column categorizes every product into a certain category (say, for eg., Electronics, sports etc.) The way the other 2 columns are supposed to be interpreted is that one product may belong to more than one category. If it doesn't, the other 2 columns simply won't contain any data. And that is the reason why those 2 columns, in particular, have missing data.

We decided to implement a simple solution for this missing data problem. The categories start from '1' and range till '18'. So, we decided to replace all the NaN values with 0. Thus, we made a new category '0' for all the missing or uncategorized data points (uncategorized in the sense that they don't belong to more than one or more than two categories).

User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase	
0	1000001	P00069042	F	0-17	10	A	2	0	3	NaN	NaN	8370
1	1000001	P00240942	F	0-17	10	A	2	0	1	6.0	14.0	15200
2	1000001	P00087842	F	0-17	10	A	2	0	12	NaN	NaN	1422
3	1000001	P00085442	F	0-17	10	A	2	0	12	14.0	NaN	1057
4	1000002	P00285442	M	55+	16	C	4+	0	8	NaN	NaN	7969

Data in its raw form

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 537577 entries, 0 to 537576
Data columns (total 12 columns):
User_ID                537577 non-null int64
Product_ID             537577 non-null object
Gender                 537577 non-null object
Age                   537577 non-null object
Occupation             537577 non-null int64
City_Category          537577 non-null object
Stay_In_Current_City_Years  537577 non-null object
Marital_Status         537577 non-null int64
Product_Category_1     537577 non-null int64
Product_Category_2     370591 non-null float64
Product_Category_3     164278 non-null float64
Purchase               537577 non-null int64
dtypes: float64(2), int64(5), object(5)
memory usage: 49.2+ MB
None

```

Metadata

```

[8 rows x 7 columns]
User_ID                0
Product_ID             0
Gender                 0
Age                   0
Occupation             0
City_Category          0
Stay_In_Current_City_Years  0
Marital_Status         0
Product_Category_1     0
Product_Category_2     166986
Product_Category_3     373299
Purchase               0

```

Missing Values

We removed 2 columns, User_ID and Product_ID as both the attributes were not significant for our analysis. Both the columns won't be that helpful when we choose to carry out a predictive analysis on our data. The reason is that, we are trying to predict the purchase amount as a part of this project. The IDs are a label, mostly for data storing, and being so distinct as they are, they'd rather make the model more specific and result in overfitting, if included at all. Thus, we decided to remove those 2 columns.

Another data inconsistency problem was the presence of '+' sign in 'Age' and 'Stay_In_Current_City_Years'. We are planning to use these bins as levels, giving them an order in their occurrence. '+' sign will introduce an inconsistency and thus we decided upon removing the '+' sign. We used the 'strip' function in combination with 'apply' function on these particular columns to get rid of the '+' signs.

```
460 #Reading the data:
461 df = pd.read_csv("BlackFriday.csv")
462 print(df.head())
463 print("-----")
464
465 #-----
466
467 #info will print the number of entries (non-null) and their data type. Describe gives statistical description:
468 print(df.info())
469 print("-----")
470 print(df.describe())
471 print("-----")
472 #-----
473
474 #fillna isn't working. Either that or taking just too long:
475 #Any other methods/suggestions to fill the 'na's?
476
477 #df=df.fillna(df.mean()['Product_Category_1':'Product_Category_2'])
478 #print(df.head())
479
480 #-----
481
482 #To see columnwise total number of missing data points:
483 print(df.isnull().sum())
484 print("-----")
485
486 df=df.replace(np.nan,0)
487 print(df.info())
488 print("-----")
489
490 print(df.isnull().sum())
491 print("-----")
492
493 #-----
494
```

Data Cleaning - 1

```

518
519 #Taking column names in a list:
520 col_names=df.columns.values.tolist()
521
522 clean_df=pd.DataFrame()
523
524 #Columns that won't contribute to machine learning:
525 to_drop=["User_ID","Product_ID"]
526 for col in col_names:
527     if col in to_drop:
528         continue
529     else:
530         clean_df[col]=df[col]
531
532 to_convert=["Product_Category_2","Product_Category_3"]
533 print("Trying to coerce PC2 and PC3 into int...")
534 for col in col_names:
535     if col in to_convert:
536         print("Data type of",col,"is",df[col].dtype)
537         clean_df[col]=clean_df[col].astype(int)
538         print("Now it is",df[col].dtype)
539
540 print("Printing head...")
541 print(clean_df.head())
542
543 #-----
544
545 to_clean=["Stay_In_Current_City_Years"]
546 col_names1=clean_df.columns.values.tolist()
547 #Not making this was screwing with the cleaning of '+' cause I was using the column names list from unclean dataframe.
548 #Cleaning the data for '+' signs. I need to put this out in the comments: Categories: '4' for Years in city and '55' for age are
549
550 col="Stay_In_Current_City_Years"
551 clean_df[col]=clean_df[col].apply(lambda x: x.strip("+")).astype(int)
552
553 col="Age"
554 clean_df[col]=clean_df[col].apply(lambda x: x.strip("+"))

```

Data Cleaning - 2

We had also gone through the Descriptive Analysis to get to know better the numerical attributes of our Dataset.

```

normaling for column: Occupation, Stay_In_Current_City_Years, Marital_Status, Product_Category_1, Product_Category_2, Product_Category_3, Purchase

```

	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
count	537577.00000	537577.00000	537577.00000	537577.00000	537577.00000	537577.00000	537577.00000
mean	8.08271	1.859458	0.408797	5.295546	9.842144	12.669840	9333.859853
std	6.52412	1.289828	0.491612	3.750701	4.223872	2.279938	4981.022133
min	0.00000	0.00000	0.00000	1.00000	2.00000	3.00000	185.00000
25%	2.00000	1.00000	0.00000	1.00000	8.00000	12.669840	5866.00000
50%	7.00000	2.00000	0.00000	5.00000	9.842144	12.669840	8062.00000
75%	14.00000	3.00000	1.00000	8.00000	14.00000	12.669840	12073.00000
max	20.00000	4.00000	1.00000	18.00000	18.00000	18.00000	23961.00000

Descriptive Analysis

For the Product Category 1,2 and 3 we first thought to analyze it by creating a new attribute Product Category. The Product Category will be calculated by using all 3 Product Categories.

$$\text{Product_Category} = \text{Product_Category_1} * 10000 + \text{Product_Category_2} * 100 + \text{Product_Category_3} * 1.$$

However, our assumptions related to Product Categories were wrong. This way of calculating Product Category won't be significant while generating models. So, we re-evaluated our method and later changed it.

We thought of using dummies as an alternative approach. We had 18 categories in our Product Category, so making 17 dummy variables wasn't feasible. That's why we decided to go with 3 product categories 1, 5 and 8 as their frequency in the data was very high. This is the method which we later used for our Product Category attribute.

1. Dummy_1=1, if Product_Category_1==1 or Product_Category_2==1 or Product_Category_3==1
2. Dummy_5=1, if Product_Category_1==5 or Product_Category_2==5 or Product_Category_3==5
3. Dummy_8=1, if Product_Category_1==8 or Product_Category_2==8 or Product_Category_3==8

```
576 clean_df["PC1"]=0
577 clean_df["PC5"]=0
578 clean_df["PC8"]=0
579
580 clean_df.loc[(clean_df.Product_Category_1==1)|(clean_df.Product_Category_2==1)|(clean_df.Product_Category_3==1),"PC1"]=1
581 clean_df.loc[(clean_df.Product_Category_1==5)|(clean_df.Product_Category_2==5)|(clean_df.Product_Category_3==5),"PC5"]=1
582 clean_df.loc[(clean_df.Product_Category_1==8)|(clean_df.Product_Category_2==8)|(clean_df.Product_Category_3==8),"PC8"]=1
583
```

Product Category Code

Challenges and Solution:

A further visual analysis showed us that 3 of the 18 product categories in our data showed a significantly high number of sales, when compared to others. But it is to be noted that the column labels are mere labels, barring any ordering whatsoever. So, it won't make any sense if we used these product category labels directly for our analysis as they don't mean anything. For all intent and purposes, we may replace 1 with apples, 2 with oranges and so on and ultimately it won't deprecate our analysis in any manner. This problem was realized a further down the road, and it required a pragmatic solution. As mentioned before, 3 categories showed higher sale than other. Those categories were '1', '5' and '8'. Thus, we decided to make dummy variables representing these 3 product categories. The approach used was that if either of the 3 product category columns had product classified as one of these 3 categories, that particular dummy will take the value 1. Being neither of these 3 categories but belonging to any other category was held as baseline, for both individual dummies and all 3 dummies as collective.

Thus, for eg., if a product belonged to category 1 and category 5, the PC1 and PC5 dummy would be 1 and PC8 dummy would be 0. If the product belonged only to category 0, then, PC1, PC5 and PC8 would all be 0.

Outliers:

We also chose to have a look at the outliers, if any. We wrote a custom function for that, which took one column at a time as an argument. It calculated the mean and standard deviation for that column and then eventually calculated the z-scores. We decided to take a z-score absolute value of 3 as a threshold for outlier and our function returned the row numbers of all the observations determined to be outliers, as an assistance in handling them. While calling the function, we made sure that the function was called only for the columns that had the data type integer or float, for only they can be evaluated for outliers.

```
558 out_idx={}
559 for col in col_names1:
560     print("Running for column",col,"Data type:",df[col].dtype)
561     if(clean_df[col].dtype=="int64" or clean_df[col].dtype=="float64"):
562         out=list(outlier_test(clean_df[col]))
563         if len(out[0])>0:
564             print("Outliers found in column",col,"Number of outliers:",len(out[0]))
565             out_idx[col]=out
566     else:
567         print("No outliers in column",col)
568
```

```
31
32 def outlier_test(inp):
33     #I am using the condiiton of absolute value of z-score>3 as the parameter of being an outlier:
34     threshold=3
35     avg_inp=np.mean(inp)
36     stdev=np.std(inp)
37     zscores=[]
38     for i in inp:
39         z=(i-avg_inp)/stdev
40         zscores.append(z)
41     return np.where(np.abs(zscores) > threshold)
42
```



```

('Running for column', 'Gender', 'Data type:', dtype('O'))
('No outliers in column', 'Gender')
('Running for column', 'Age', 'Data type:', dtype('O'))
('No outliers in column', 'Age')
('Running for column', 'Occupation', 'Data type:', dtype('int64'))
('Running for column', 'City_Category', 'Data type:', dtype('O'))
('No outliers in column', 'City_Category')
('Running for column', 'Stay_In_Current_City_Years', 'Data type:', dtype('O'))
('Running for column', 'Marital_Status', 'Data type:', dtype('int64'))
('Running for column', 'Product_Category_1', 'Data type:', dtype('int64'))
('Outliers found in column', 'Product_Category_1', 'Number of outliers:', 3642)
('Running for column', 'Product_Category_2', 'Data type:', dtype('float64'))
('Running for column', 'Product_Category_3', 'Data type:', dtype('float64'))
('Running for column', 'Purchase', 'Data type:', dtype('int64'))

```

It shows outliers only for the column 'Product_Category_1'. Ironically, it won't count to anything, as discussed earlier, they are all labels and they don't have a fixed order amongst the product categories. All the other columns are outlier free. (P.S.: It shows running for each column, but only numeric columns get the 2nd output, i.e. either no outliers found, or number of outliers found. The non-numeric ones don't call the function. Corroborated from the snippet of code in image 4 and image 5)

Thus finally, the cleaned data looks as follows:

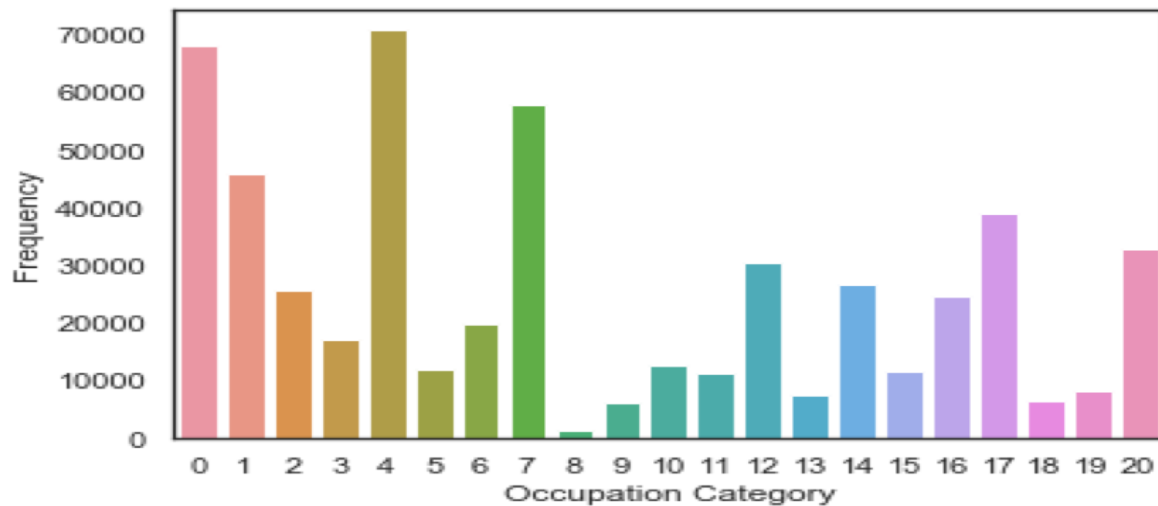
	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase	PC1	PC5	PC8
0	F	0-17	10	A	2	0	3	0	0	8370	0	0	0
1	F	0-17	10	A	2	0	1	6	14	15200	1	0	0
2	F	0-17	10	A	2	0	12	0	0	1422	0	0	0
3	F	0-17	10	A	2	0	12	14	0	1057	0	0	0
4	M	55	16	C	4	0	8	0	0	7969	0	0	1
5	M	26-35	15	A	3	0	1	2	0	15227	1	0	0
6	M	46-50	7	B	2	1	1	8	17	19215	1	0	1
7	M	46-50	7	B	2	1	1	15	0	15854	1	0	0
8	M	46-50	7	B	2	1	1	16	0	15686	1	0	0
9	M	26-35	20	A	1	1	8	0	0	7871	0	0	1
10	M	26-35	20	A	1	1	5	11	0	5254	0	1	0
11	M	26-35	20	A	1	1	8	0	0	3957	0	0	1
12	M	26-35	20	A	1	1	8	0	0	6873	0	0	1
13	M	26-35	20	A	1	1	1	2	5	15665	1	1	0
14	F	51-55	9	A	1	0	5	8	14	5378	0	1	1
15	F	51-55	9	A	1	0	4	5	0	2079	0	1	0
16	F	51-55	9	A	1	0	2	3	4	13855	0	0	0
17	F	51-55	9	A	1	0	5	14	0	8851	0	1	0
18	M	36-45	1	B	1	1	1	14	16	11788	1	0	0
19	M	26-35	12	C	4	1	1	5	15	19614	1	1	0

Another point to be noted is that, although this is termed to be clean data, this is still raw data when it comes to be used for analysis. Every analysis and predictive modeling method will require more cleaning and munging so that the data is fit and ready for that particular predictive model.

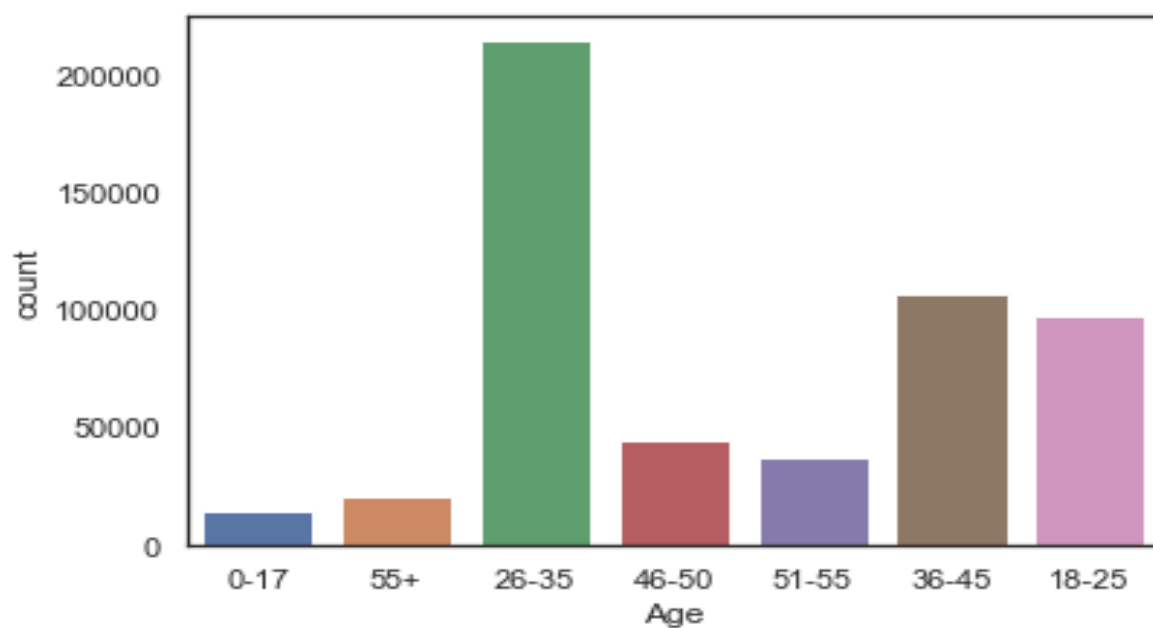
Descriptive Statistics:

We created visualization graphs of significant attributes with their frequency to better visualize our data.

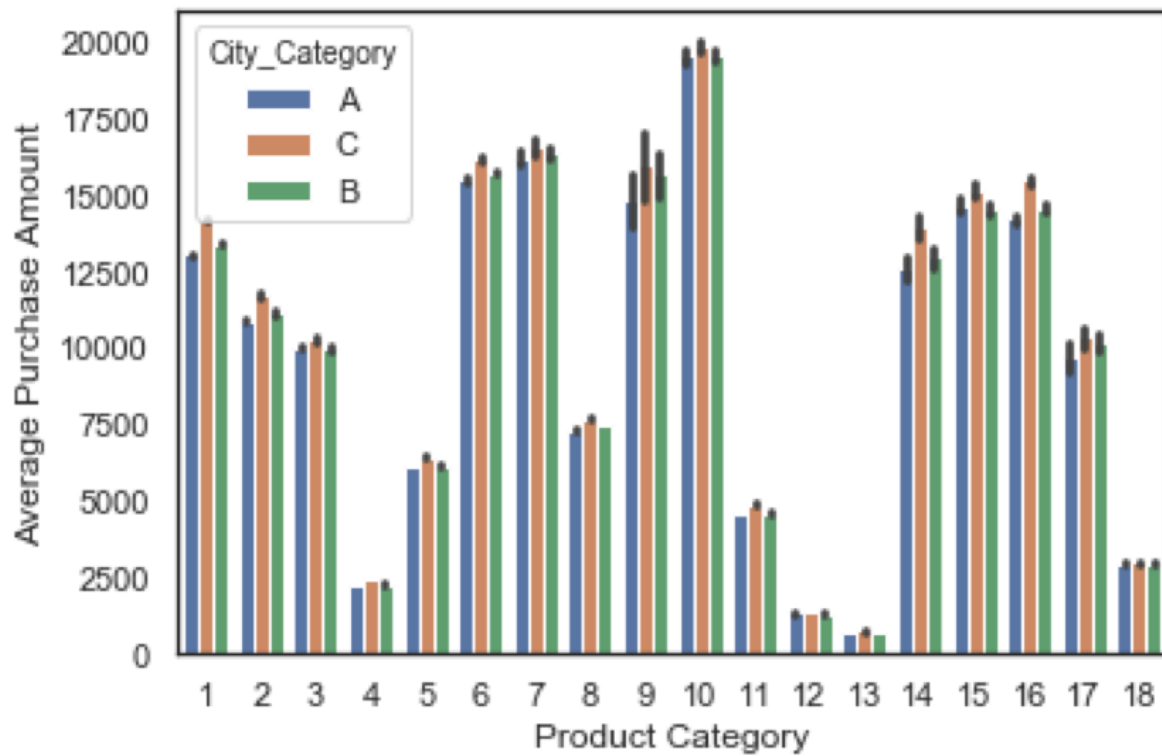
1. Occupation – Below is the Occupation of people with respect to their purchase count.
We found out that the greatest number of purchase count were by occupation 0, 1, 4, 7, 17 and 20.



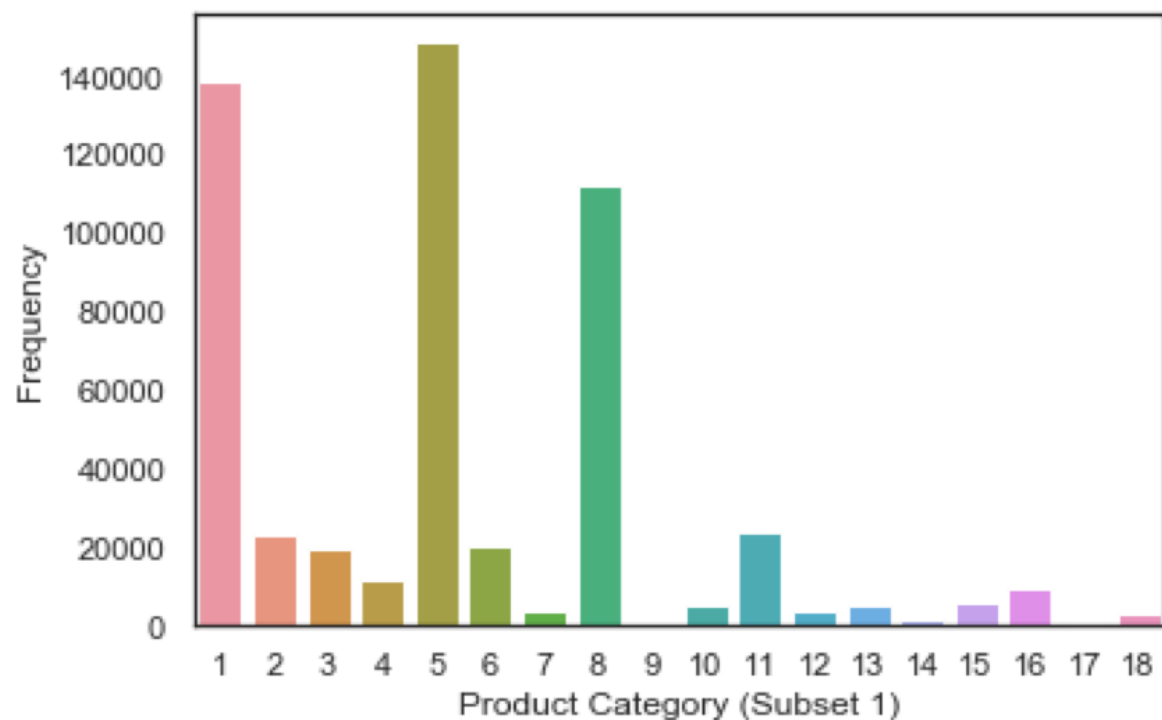
2. Age – Age Group of people and their respective count of purchase. The age groups between 18-45 are the ones who purchases the most.



3. Product & City – Here is the product and city category with respect to their Purchase Amount.



4. Product Category – Below is the Product Category 1 and its Frequency. From this we get to know that Product Category 1, 5 and 8 were purchased the most.



Linear Regression:

For generating Linear Regression model, we have created dummy variables. The dummy variables which we have created are:

1. For 'Gender': We took 'F' as our baseline, i.e. 0 and 'M' as 1.
2. For 'Age': We had 7 bins originally, so we made 6 dummy variables and 0-17 bin is the baseline (all dummies 0)
3. For 'Occupation': We selected 6 categories out of 20 that had highest number of observations. The selected categories are 0,1,4,7,17,20. All the other categories are the baseline.
4. For 'City Category': We took C as our baseline and 2 dummy variables for category A and B.
5. For 'Stay in current city', 'Marital status' and dummies for product categories 1,5 and 8 are taken as it is.

	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase	PC1	PC5	PC8
0	F	0-17	10	A	2	0	3	0	0	8378	0	0	0
1	F	0-17	10	A	2	0	1	6	14	15208	1	0	0
2	F	0-17	10	A	2	0	12	0	0	1422	0	0	0
3	F	0-17	10	A	2	0	12	14	0	1857	0	0	0
4	M	55	16	C	4	0	8	0	0	7969	0	0	1
5	M	26-35	15	A	3	0	1	2	0	15227	1	0	0
6	M	46-50	7	B	2	1	1	8	17	19215	1	0	1
7	M	46-50	7	B	2	1	1	15	0	15854	1	0	0
8	M	46-50	7	B	2	1	1	16	0	15686	1	0	0
9	M	26-35	20	A	1	1	8	0	0	7871	0	0	1
10	M	26-35	20	A	1	1	5	11	0	5254	0	1	0
11	M	26-35	20	A	1	1	8	0	0	3957	0	0	1
12	M	26-35	20	A	1	1	8	0	0	6873	0	0	1
13	M	26-35	20	A	1	1	1	2	5	15665	1	1	0
14	F	51-55	9	A	1	0	5	0	14	5378	0	1	1
15	F	51-55	9	A	1	0	4	5	0	2879	0	1	0
16	F	51-55	9	A	1	0	2	3	4	13855	0	0	0
17	F	51-55	9	A	1	0	5	14	0	8851	0	1	0
18	M	36-45	1	B	1	1	1	14	16	11788	1	0	0
19	M	26-35	12	C	4	1	1	5	15	19614	1	1	0

Original Data

	Purchase	Gender	dum_18-25	dum_26-35	dum_36-45	dum_46-50	dum_51-55	dum_55+	dum_0	dum_1	dum_4	dum_7	dum_17	dum_20	City_A	City_B	In_City	Marital_Status	P1_dum	P5_dum	P8_dum
0	8378	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	2	0	0	0	0
1	15208	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	2	0	1	0	0
2	1422	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	2	0	0	0	0
3	1857	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	2	0	0	0	0
4	7969	1	0	0	0	0	0	1	0	0	0	0	0	0	0	0	4	0	1	0	1
5	15227	1	0	1	0	0	0	0	0	0	0	0	0	0	1	0	3	0	1	0	0
6	19215	1	0	0	0	1	0	0	0	0	0	1	0	0	0	1	2	1	1	0	1
7	15854	1	0	0	0	1	0	0	0	0	0	1	0	0	0	1	2	1	1	0	0
8	15686	1	0	0	0	1	0	0	0	0	0	1	0	0	0	1	2	1	1	0	0
9	7871	1	0	1	0	0	0	0	0	0	0	0	0	1	1	0	1	1	0	0	1
10	5254	1	0	1	0	0	0	0	0	0	0	0	0	1	1	0	1	1	0	1	0
11	3957	1	0	1	0	0	0	0	0	0	0	0	0	1	1	0	1	1	0	0	1
12	6873	1	0	1	0	0	0	0	0	0	0	0	0	1	1	0	1	1	0	0	1
13	15665	1	0	1	0	0	0	0	0	0	0	0	0	1	1	0	1	1	1	1	0
14	5378	0	0	0	0	0	1	0	0	0	0	0	0	0	1	0	1	0	0	1	1
15	2879	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	1	0	0	1	0
16	13855	0	0	0	0	0	1	0	0	0	0	0	0	0	1	0	1	0	0	0	0
17	8851	0	0	0	0	0	1	0	0	0	0	0	0	0	1	0	1	0	0	1	0
18	11788	1	0	0	1	0	0	0	0	1	0	0	0	0	1	1	1	1	1	0	0
19	19614	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	4	1	1	1	0

Data with Dummy Variables

```

=====
                        OLS Regression Results
=====
Dep. Variable:          y          R-squared:          0.300
Model:                  OLS        Adj. R-squared:       0.300
Method:                  Least Squares   F-statistic:       9689.
Date:                    Tue, 02 Apr 2019   Prob (F-statistic): 0.00
Time:                    13:54:31    Log-Likelihood:    -4.1951e+06
No. Observations:        430061    AIC:               8.390e+06
Df Residuals:            430041    BIC:               8.390e+06
Df Model:                 19
Covariance Type:         nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	8920.7110	42.033	212.234	0.000	8838.329	9003.093
x1	79.1526	15.015	5.272	0.000	49.723	108.582
x2	65.2426	42.867	1.522	0.128	-18.775	149.260
x3	204.4134	40.564	5.039	0.000	124.908	283.918
x4	281.7199	41.733	6.751	0.000	199.925	363.515
x5	236.9423	45.814	5.172	0.000	147.148	326.737
x6	505.6492	46.677	10.833	0.000	414.165	597.134
x7	320.4790	51.135	6.267	0.000	220.256	420.701
x8	-161.1987	20.535	-7.850	0.000	-201.447	-120.951
x9	-182.7126	24.047	-7.598	0.000	-229.844	-135.581
x10	24.0936	22.931	1.051	0.293	-20.851	69.038
x11	-44.1248	22.024	-2.004	0.045	-87.290	-0.959
x12	76.5931	25.807	2.968	0.003	26.013	127.174
x13	-138.8301	27.826	-4.989	0.000	-193.368	-84.292
x14	-631.7454	17.163	-36.808	0.000	-665.384	-598.106
x15	-471.0299	15.221	-30.946	0.000	-500.862	-441.197
x16	6.9117	4.940	1.399	0.162	-2.771	16.594
x17	-37.8679	13.791	-2.746	0.006	-64.898	-10.837
x18	5096.3202	15.186	335.600	0.000	5066.557	5126.084
x19	-2158.8412	13.822	-156.184	0.000	-2185.933	-2131.750

```

=====
Omnibus:                24688.718    Durbin-Watson:         2.001
Prob(Omnibus):           0.000        Jarque-Bera (JB):      34278.762
Skew:                    0.526        Prob(JB):              0.00
Kurtosis:                 3.899        Cond. No.              44.1
=====

```

Output

Results:

The take outs from Regression Model are -

1. There were 3 explanatory variables having a p-value higher than 0.05, thus showing they are less significant.
2. The value of R-squared and adjusted R-squared was 0.3.
3. For this data set, males tend to spend more than females.
4. Regression shows that every additional year in the city leads to a decrease in the amount of purchase.
5. Spending pattern over age is not plainly increasing or decreasing.

Random Forest Regressor:

The random forest regressor model will make a forest out of a certain number of decision trees. The decision trees are made by pulling out random samples of data from input data. The size of sample stays the same, but bootstrapping enables the model to pull different samples for making different trees.

The trees have a specified number of maximum leaf nodes. These are the maximum number of data points to be assigned to one leaf node. The data is classified into trees by finding patterns in the data. Data points showcasing similar patterns in explanatory and target variable end up being assigned to the same leaf node. What random forest regressor will do is that when it is used to finally predict for an unknown, never seen before data record, it will start traversing from the root node and follow the branches on the basis of the value of the explanatory variables from the data entry that is used for prediction. It will eventually reach one of the leaf nodes based on the values of these variables, thus classifying the data entry into a group showcasing similar patterns. Eventually, an average of the values of target variables for all the data points assigned to that particular leaf node is given to the new data point as a result of prediction.

We wrote a function for random forest regressor that took a data frame as its input. The data frame from earlier section (cleaned data) is passed to this function.

```
48 def forest_regressor(for_for):
49
50     #-----
51
52     start=time.time()
53
54     #-----
55
56     x_for=for_for.iloc[:,0:9].values
57     y_for=for_for.iloc[:,9].values
58
59     x_for_train,x_for_test,y_for_train,y_for_test=train_test_split(x_for,y_for,test_size=0.2,random_state=0)
60     print(x_for_train[:5])
61     print(y_for_train[:5])
62
63     #-----
64
65     #Let's start with, encoding:
66
67     lab_encoder=LabelEncoder()
68     #encode=["Gender","City_Category"]
69     #for col in encode:
70     #For Gender:
71     x_for_train[:,0]=lab_encoder.fit_transform(x_for_train[:,0])
72     x_for_test[:,0]=lab_encoder.fit_transform(x_for_test[:,0])
73
74     #For Age Groups:
75     x_for_train[:,1]=lab_encoder.fit_transform(x_for_train[:,1])
76     x_for_test[:,1]=lab_encoder.fit_transform(x_for_test[:,1])
77
78     #For City Category:
79     x_for_train[:,3]=lab_encoder.fit_transform(x_for_train[:,3])
80     x_for_test[:,3]=lab_encoder.fit_transform(x_for_test[:,3])
81
82     print(x_for_train[:5])
```

```

86     #Now we gonna do scaling:
87     scalar=StandardScaler()
88     x_for_train=scalar.fit_transform(x_for_train)
89
90     test_scalar=StandardScaler()
91     x_for_test=test_scalar.fit_transform(x_for_test)
92
93     print(x_for_train[:5])
94
95     #-----
96
97     end=time.time()
98     print("-----")
99     print("Time taken for pre-processing for the Random Forest Regressor is:",end-start,"seconds")
100    print("-----")
101
102    #-----
103
104    max_leaves=input("Maximum number of leaf nodes you wish to have: ")
105    print("-----")
106    max_leaves=int(max_leaves)
107

```

```

113
114    model=RandomForestRegressor(max_leaf_nodes=max_leaves,random_state=0)
115    model.fit(x_for_train,y_for_train)
116    pred=model.predict(x_for_test)
117

```

```

124    #
125
126    mae=mean_absolute_error(y_for_test,pred)
127    print("For",max_leaves,"nodes, the mean absolute error:",mae)
128    print("-----")
129    mse=mean_squared_error(y_for_test,pred)
130    rmse=np.sqrt(mse)
131    print("Root Mean Squared value for",max_leaves,"nodes is:",rmse)
132    print("-----")
133

```

The following snippet of codes show the pre-processing involved for the random forest regressor, the implementation of the model and the calculation of metrics involved in evaluating the accuracy of the model.

The first step is to separate the input data into explanatory and target variables, done here by splitting the data into 2 numpy arrays, 'x' and 'y'. Then we use the train_test_split function from sklearn package's model_selection class to split these 2 arrays into 2 subsections, one for training the model and another for testing its performance. We chose to split the data 80:20, i.e. 80% training and 20% testing. We are choosing the data not to be sampled randomly.

Here's a look at the raw split training data (both explanatory and target):

```
[[ 'M' '26-35' 0 'B' 3 0 8 0 0]
 [ 'M' '26-35' 2 'B' 1 0 5 0 0]
 [ 'M' '26-35' 0 'A' 3 1 8 14 0]
 [ 'F' '46-50' 1 'C' 1 1 3 4 12]
 [ 'F' '26-35' 3 'B' 1 1 5 14 0]]
[ 7837 8762 7778 13701 7158]
```

We then used LabelEncoder so that the categories in our data get encoded into simpler labels. LabelEncoder is a function from sklearn package's preprocessing class. The simpler labels are general numbers starting from 0 and so forth. The labels are assigned as per the comparison of models. For eg, for age, 0-17 will be assigned as 0, since the starting of this bin is 0 (lowest bin) and so on. We used LabelEncoder on 3 columns, 0,1 and 3 which are gender, age and city category. The outcome of label encoding is as follows:

```
[[ 1 2 0 1 3 0 8 0 0]
 [ 1 2 2 1 1 0 5 0 0]
 [ 1 2 0 0 3 1 8 14 0]
 [ 0 4 1 2 1 1 3 4 12]
 [ 0 2 3 1 1 1 5 14 0]]
```

Finally, we decided to scale the data. Although, there is not a drastic change in the range of data, the 12s and 14s (and similar data), may drive the decisions comparatively more than the other variables. So, we have used the StandardScaler from the sklearn's preprocessing class in order to scale the data to have mean 0 and standard deviation 1, for every column individually, thus maintaining the data's characteristics of variance. The output is as follows:

```
[[ 0.57015696 -0.36616352 -1.23968144 -0.05494249 0.88240648 -0.83107805
 0.7225614 -1.09294693 -0.61819295]
 [ 0.57015696 -0.36616352 -0.93292539 -0.05494249 -0.66780477 -0.83107805
 -0.07724542 -1.09294693 -0.61819295]
 [ 0.57015696 -0.36616352 -1.23968144 -1.37130026 0.88240648 1.20325642
 0.7225614 1.16032749 -0.61819295]
 [-1.75390301 1.1124876 -1.08630342 1.26141528 -0.66780477 1.20325642
 -0.61044997 -0.44915424 1.29674464]
 [-1.75390301 -0.36616352 -0.77954737 -0.05494249 -0.66780477 1.20325642
 -0.07724542 1.16032749 -0.61819295]]
```

Finally, we make an object instance of Random Forest regressor and we use the training explanatory and target arrays to fit the model. Then we use the testing explanatory array to make predictions. We resorted to evaluating 3 types of forecast errors; mean absolute error,

mean squared error and root mean squared error. We chose to make our model more user driven so that we can see the change in performance by tuning a few parameters. Thus, we made a functionality where the user can enter the maximum number of leaf nodes he or she wants. By default, the algorithm will make a forest with 100 trees (that's the default value of `n_estimators`). Bootstrapping is used as a default of the algorithm and we didn't see any point in not having bootstrapping, as it will help create a more generalized model due to the randomness. The default criterion is 'mse' (mean squared error) and the algorithm will try making a model that minimizes the mean squared error. The output in pictures is for a forest with all the trees having 1250 number of maximum leaf nodes:

```
-----
('Time taken for pre-processing for the Random Forest Regressor is:', 1.7082159519195557, 'seconds')
-----
Maximum number of leaf nodes you wish to have: 1250
-----
/anaconda3/lib/python2.7/site-packages/sklearn/ensemble/forest.py:246: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)
('Time taken for Random Forest Regressor to execute:', 10.429543018341064, 'seconds.')
-----
('For', 1250, 'nodes, the mean absolute error:', 2186.0758406829414)
-----
('Root Mean Squared value for', 1250, 'nodes is:', 2904.635055088412)
-----
```

Challenges and tweaked approach:

No major challenges were faced when carrying out the building and execution of this model. Reason behind this is the code was pretty straight forward. Understanding the concept and thinking of the data munging was a challenge. We tried making the code as dynamic as possible, but we couldn't make the code generalized enough for encoding; we had to hard code the column names and so forth.

Looking back, we think that we may have been able to create a better model had we used some different components of data. We interpreted this as a forest of decision trees making decision on the basis of patterns in the explanatory variables and hence, despite knowing that the labels in product categories are not ordered, we chose to use them as they still identify product categories distinctly. But we think that there is a possibility that rather than using the product category columns, the model would have performed better, had we used the dummy variables we made for the linear regression model.

Random Forest Classifier:

For most of the parts, random forest classifier works same as the random forest regressor. The modeling technique will make a forest of a certain number of trees that would map the pattern in the explanatory variables. The broad difference is that rather than having an average of the target variable as the value of the node, the classifier requires binning or discretizing the numeric target variable and eventually, it will classify data points on the basis of the explanatory variables and a given number of maximum data points will be assigned to a leaf node. The mode of the bins of these data points is assigned as the resultant category for this leaf node.

We strived and succeeded in making the model user driven. The system will ask for number of bins the user wishes to discretize the 'Purchase' column into. Based on the number entered, the system will calculate the bin width by using the maximum and minimum value from the target column. After getting the width and the number of bins, we make 2 lists, one with bin start value and one with respective bin label. Then we used the 'cut' function from pandas which will take the array or data-frame to be discretized, the bin start point list and the bin label list. We then used the LabelEncoder on the same column as we used it on in the regressor model, i.e. 0,1 and 3; gender, age and city category respectively. Finally, we used Standard Scaler on the data.

Running a Random Forest Classifier requires us to initialize model instance. We then fit the model on the training data (both explanatory and discretized target). Finally, we will use the trained model and testing explanatory data to make predictions. In order to evaluate the model's accuracy we used 3 parameters; confusion matrix, classification report and accuracy score. All these are implemented using confusion_matrix, classification_report, accuracy_score functions from sklearn package's metrics class.

A bit about the Random Forest Classifier function: We have made the maximum number of leaf nodes parameter to be user driven so that the user can actually see the variation in performance of model when this and number of bins parameter is tuned. We used a different number of trees parameter (n_estimators) than the default one. We fixed it to be 150. However, we can also make it user driven, it is a matter of 2 lines of code. The criterion used is 'Gini Index'. Another available criterion was 'information gain' (Entropy). Both are a measure of impurities. Entropy will be 0 if all the nodes belong to 1 category (Entropy is the change or

variance) and will be the highest of all the nodes are evenly distributed over different categories. Exact same goes for Gini Index. Gini Index is faster, due to the mathematical ease of computation. Both give more or less the same output.

$$Entropy = - \sum_j p_j \log_2 p_j$$

$$Gini = 1 - \sum_j p_j^2$$

```

139 def forest_classifier(for_for):
140
141     #-----
142
143     #We gotta make bins for purchase amount:
144     print("-----")
145     num_bins=input("Enter the number of bins: ")
146     print("-----")
147     num_bins=int(num_bins)
148
149     #-----
150
151     start=time.time()
152
153     #-----
154
155     bin_width=((max(for_for.Purchase))-(min(for_for.Purchase)))/num_bins
156     print("For",num_bins,"bins, bin width=",bin_width)
157     print("-----")
158
159     bins=[]
160     lab=[]
161     i=0
162     while i<num_bins+1:
163         if i==0:
164             bins.append(min(for_for.Purchase))
165             lab.append(i)
166         else:
167             bins.append(bins[i-1]+bin_width)
168             lab.append(i)
169         i=i+1
170
171     #print(bins)
172     lab=lab[:-1]
173     #print(lab)
174

```

```

177 x_for=for_for.iloc[:,0:9].values
178 y_for=for_for.iloc[:,9].values
179
180 x_for_train,x_for_test,y_for_train,y_for_test=train_test_split(x_for,y_for,test_size=0.2,random_state=0)
181
182 #-----
183
184 y_train_bin=pd.cut(y_for_train,bins=bins,labels=lab)
185 print(y_train_bin[:20])
186 print("-----")
187
188 y_test_bin=pd.cut(y_for_test,bins=bins,labels=lab)

```

```

197 lab_encoder=LabelEncoder()
198 #encode=["Gender","City_Category"]
199 #for col in encode:
200 #For Gender:
201 x_for_train[:,0]=lab_encoder.fit_transform(x_for_train[:,0])
202 x_for_test[:,0]=lab_encoder.fit_transform(x_for_test[:,0])
203
204 #For Age Groups:
205 x_for_train[:,1]=lab_encoder.fit_transform(x_for_train[:,1])
206 x_for_test[:,1]=lab_encoder.fit_transform(x_for_test[:,1])
207
208 #For City Category:
209 x_for_train[:,3]=lab_encoder.fit_transform(x_for_train[:,3])
210 x_for_test[:,3]=lab_encoder.fit_transform(x_for_test[:,3])
211
212 #-----
213
214 #Now we gonna do scaling:
215 scalar=StandardScaler()
216 x_for_train=scalar.fit_transform(x_for_train)
217
218 test_scalar=StandardScaler()
219 x_for_test=test_scalar.fit_transform(x_for_test)
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239 classifier=RandomForestClassifier(max_leaf_nodes=max_leaves,bootstrap=True,criterion='gini',n_estimators=150)
240 classifier.fit(x_for_train,y_train_bin)
241 pred=classifier.predict(x_for_test)
242
243 #-----

```

```

250
251 print("The confusion matrix for the Random Forest Classifier is as follows:")
252 print(confusion_matrix(y_test_bin,pred))
253 print("-----")
254 print("The classification report for the Random Forest Classifier is as follows:")
255 print(classification_report(y_test_bin,pred))
256 print("-----")
257 print("Accuracy for the Random Forest Classifier model:",(accuracy_score(y_test_bin,pred))*100,"%")
258 print("-----")
259

```

Output:

The output is for 5 bins, 1500 maximum number of leaf nodes, 150 trees in forest and bootstrapped model:

```

-----
Enter the number of bins: 5
-----
('For', 5, 'bins, bin width=', 4755)
-----
[1, 1, 1, 2, 1, ..., 3, 1, 2, 2, 4]
Length: 20
Categories (5, int64): [0 < 1 < 2 < 3 < 4]
-----

```

```
Maximum number of leaf nodes you wish to have: 1500
```

```
('Time taken for execution of the Random Forest Classifier is:', 88.80100893974304, 'seconds')
```

```
The confusion matrix for the Random Forest Classifier is as follows:
```

```
[[ 7743  7631   703  1704    26]
 [ 1945 38905  2721  3547    66]
 [   90  6182  5930  8272   295]
 [   58   262  1271 13255   650]
 [   36   167   142  5023   892]]
```

```
The classification report for the Random Forest Classifier is as follows:
```

	precision	recall	f1-score	support
0	0.78	0.43	0.56	17807
1	0.73	0.82	0.78	47184
2	0.55	0.29	0.38	20769
3	0.42	0.86	0.56	15496
4	0.46	0.14	0.22	6260
micro avg	0.62	0.62	0.62	107516
macro avg	0.59	0.51	0.50	107516
weighted avg	0.64	0.62	0.60	107516

```
('Accuracy for the Random Forest Classifier model:', 62.060530525689195, '%')
```

The accuracy obtained is 62.06%. However, the highest accuracy we got was 62.18% for the same setting, except that we had ‘information gain’ criterion instead of ‘Gini Index’ criterion.

Challenges and setbacks:

This model was not only the best in all the models we implemented but was also the most complicated and challenging one. We faced a couple of challenges that required some obvious yet oblivious solutions. Initially, we tried fitting the model by making 20 bins. This was when everything was still in development and we had to hard code the number of models. Thus, we couldn’t run multiple iterations for different number of models. But even before we could successfully execute the model for 20 bins, we were faced with an error:

ValueError: Input contains NaN, infinity or a value too large for dtype('float64').

We couldn’t get the complete meaning of the error at first. After some thinking, we found out that the binning was faulty. So we had a look at the bins and the maximum number in the ‘Purchase’ column:

```
count    Purchase
mean      537577.000000
std       9333.859853
min       4981.022133
25%       185.000000
50%       5866.000000
75%       8062.000000
max      12073.000000
max      23961.000000
```

The problem was that the bin width was 1188.8. There were a few data points that were higher than the closing value of the last bin, by only a slight decimal margin. As a result, the 'cut' function couldn't put them in a bin and thus labeled them to have a bin value NaN. And since the Random Forest Classifier can't handle NaN values in the target variable, we were getting these errors.

	counts	freqs		counts	freqs
categories			categories		
0.0	1628	0.015142	0.0	6484	0.015077
1.0	4807	0.044710	1.0	19252	0.044766
2.0	5754	0.053518	2.0	23159	0.053851
3.0	5618	0.052253	3.0	22706	0.052797
4.0	13788	0.128241	4.0	54690	0.127168
5.0	12377	0.115118	5.0	49254	0.114528
6.0	14026	0.130455	6.0	55555	0.129179
7.0	6993	0.065041	7.0	27870	0.064805
8.0	8009	0.074491	8.0	31882	0.074134
9.0	7672	0.071357	9.0	31616	0.073515
10.0	3815	0.035483	10.0	15118	0.035153
11.0	1273	0.011840	11.0	4774	0.011101
12.0	5928	0.055136	12.0	24138	0.056127
13.0	6830	0.063525	13.0	27032	0.062856
14.0	432	0.004018	14.0	1666	0.003874
15.0	2306	0.021448	15.0	9033	0.021004
16.0	4555	0.042366	16.0	18931	0.044019
17.0	1218	0.011329	17.0	5068	0.011784
18.0	1	0.000009	18.0	9	0.000021
19.0	484	0.004502	19.0	1819	0.004230
NaN	2	0.000019	NaN	5	0.000012
	counts	freqs			

So, we just wrote 2 lines of code that filled the NaN values (using fillna function) and the NaN were replaced with the last bin label, thus coercing them into the last bin.

```

189
190     y_train_bin=y_train_bin.fillna(num_bins-1)
191     y_test_bin=y_test_bin.fillna(num_bins-1)
192

```

Then we could finally execute the model. After execution, we faced another issue. The issue was low accuracy. With a setting in which we made a forest with 150 trees, 1500 maximum

leaf nodes and Gini Index criterion for the data discretized into 20 bins, we got an accuracy of 34.49%.

Eventually, we arrived at the conclusion that 20 bins were resulting in too specific of a model and hence the overfitting impeded the model from having a better accuracy. Thus we decided to reduce down the number of bins and did the same thing (hard coding) for 5 bins rather than 20 and our accuracy went up to what was achieved.

Naïve Bayes Model:

Naïve Bayes algorithm uses Bayes theorem of probability to predict the class (or in our case bin) of an unknown data set. The base assumption in Naïve Bayes method is that all the predictors or the explanatory variables are independent of each other. Naïve Bayes is very simple to implement and a quick predictive technique. The algorithm focuses on the frequency and the resulting conditional probability and likelihood of a data point belonging to a certain bin.

Naïve Bayes required the exact same pre-processing of data as Random Forest Classifier. We made Naïve Bayes as user driven as the Random Forest Classifier. The user is asked to enter the number of bins to be used to discretize the numeric 'Purchase' column. Then we calculate the bin width and using that we make 2 lists for bin start values and bin labels. Then we use the 'cut' function to discretize the column. We then use 'fillna' to replace the NaN values (result of slightly larger data values being binned as NaN) by the label of the last bin. Then we split data into training and testing data (80:20). We make a copy of the explanatory data set before scaling.

Since we are using Gaussian Naïve Bayes, which requires data to be continuous, we'd be scaling the data. Before scaling, we'd use LabelEncoder on the columns 0,1 and 3; gender, age and city category to label them. We are using GaussianNB from sklearn package's naïve_bayes class. We create an instance for the model and fit the model onto the labeled and scaled explanatory training data and corresponding binned testing data. Then we use the same performance metrics as Random Forest Classifier to evaluate the performance of the model.

```

276     #We gotta make bins for purchase amount:
277     print("-----")
278     num_bins=input("Enter the number of bins: ")
279     print("-----")
280     num_bins=int(num_bins)
281
282     bin_width=((max(for_for.Purchase))-(min(for_for.Purchase)))/num_bins)
283     print("For",num_bins,"bins, bin width=",bin_width)
284     print("-----")
285
286     bins=[]
287     lab=[]
288     i=0
289     while i<num_bins+1:
290         if i==0:
291             bins.append(min(for_for.Purchase))
292             lab.append(i)
293         else:
294             bins.append(bins[i-1]+bin_width)
295             lab.append(i)
296         i=i+1
297
298     #print(bins)
299     lab=lab[:-1]
300     #print(lab)
301
302     #-----
303
304     x_for=for_for.iloc[:,0:9].values
305     y_for=for_for.iloc[:,9].values
306
307     x_for_train,x_for_test,y_for_train,y_for_test=train_test_split(x_for,y_for,test_size=0.2,random_state=0)
308
309     x_for_train1=x_for_train
310     x_for_test1=x_for_test
311     #-----

```

```

312
313     y_train_bin=pd.cut(y_for_train,bins=bins,labels=lab)
314     print(y_train_bin[:20])
315     print("-----")
316
317     y_test_bin=pd.cut(y_for_test,bins=bins,labels=lab)
318
319     y_train_bin=y_train_bin.fillna(num_bins-1)
320     y_test_bin=y_test_bin.fillna(num_bins-1)
321
322     #-----
323
324     #Let's start with, encoding:
325
326     lab_encoder=LabelEncoder()
327     #encode=["Gender","City_Category"]
328     #for col in encode:
329     #For Gender:
330     x_for_train[:,0]=lab_encoder.fit_transform(x_for_train[:,0])
331     x_for_test[:,0]=lab_encoder.fit_transform(x_for_test[:,0])
332
333     #For Age Groups:
334     x_for_train[:,1]=lab_encoder.fit_transform(x_for_train[:,1])
335     x_for_test[:,1]=lab_encoder.fit_transform(x_for_test[:,1])
336
337     #For City Category:
338     x_for_train[:,3]=lab_encoder.fit_transform(x_for_train[:,3])
339     x_for_test[:,3]=lab_encoder.fit_transform(x_for_test[:,3])
340
341     #-----

```



```

343     #Now we gonna do scaling:
344     scalar=StandardScaler()
345     x_for_train=scalar.fit_transform(x_for_train)
346
347     test_scalar=StandardScaler()
348     x_for_test=test_scalar.fit_transform(x_for_test)
349
350     #-----

```

```

357
358     gnb=GaussianNB()
359     gnb.fit(x_for_train,y_train_bin)
360     pred=gnb.predict(x_for_test)
361
362     #-----
363

```

```

368     #
369
370     print("The confusion matrix for the Gaussian Naive Bayes model is as follows:")
371     print(confusion_matrix(y_test_bin,pred))
372     print("-----")
373     print("The classification report for the Gaussian Naive Bayes model is as follows:")
374     print(classification_report(y_test_bin,pred))
375     print("-----")
376     print("Accuracy for the Gaussian Naive Bayes model:",(accuracy_score(y_test_bin,pred))*100,"%")
377     print("-----")

```

As we mentioned during the final presentation, we thought that multinomial Naïve Bayes may be a better model for our data and much of it is categorical data and multinomial Naïve Bayes fits better on such data. We included multinomial Naïve Bayes functionality under the Gaussian Naïve Bayes method, using the unscaled and unlabeled data set. However, Gaussian Naïve Bayes proved to be performing better, so our original work holds better of the new work. Both the models, Gaussian and Multinomial, by default, carry out Laplace smoothing on the data.

```

378
379     mnb=MultinomialNB()
380     mnb.fit(x_for_train1,y_train_bin)
381     pred1=mnb.predict(x_for_test1)
382
383     print("The confusion matrix for the Multinomial Naive Bayes model is as follows:")
384     print(confusion_matrix(y_test_bin,pred1))
385     print("-----")
386     print("The classification report for the Multinomial Naive Bayes model is as follows:")
387     print(classification_report(y_test_bin,pred1))
388     print("-----")
389     print("Accuracy for the Multinomial Naive Bayes model:",(accuracy_score(y_test_bin,pred1))*100,"%")
390     print("-----")
391

```

Output of the Naïve Bayes predictor (both) for 3 bins:

```
-----  
Enter the number of bins: 3  
-----
```

```
('For', 3, 'bins, bin width=', 7925)  
-----
```

```
[0, 1, 0, 1, 0, ..., 1, 0, 1, 1, 2]  
Length: 20  
Categories (3, int64): [0 < 1 < 2]  
-----
```

```
-----  
The confusion matrix for the Gaussian Naive Bayes model is as follows:
```

```
[[43712 10364 1262]  
 [16143 23820 303]  
 [ 4309  7157  446]]  
-----
```

```
The classification report for the Gaussian Naive Bayes model is as follows:
```

	precision	recall	f1-score	support
0	0.68	0.79	0.73	55338
1	0.58	0.59	0.58	40266
2	0.22	0.04	0.06	11912
micro avg	0.63	0.63	0.63	107516
macro avg	0.49	0.47	0.46	107516
weighted avg	0.59	0.63	0.60	107516

```
-----
```

```
('Accuracy for the Gaussian Naive Bayes model:', 63.22593846497265, '%')  
-----
```

```
The confusion matrix for the Multinomial Naive Bayes model is as follows:
```

```
[[45818 8381 1139]  
 [22604 17262 400]  
 [ 6349  5308  255]]  
-----
```

```
The classification report for the Multinomial Naive Bayes model is as follows:
```

	precision	recall	f1-score	support
0	0.61	0.83	0.70	55338
1	0.56	0.43	0.48	40266
2	0.14	0.02	0.04	11912
micro avg	0.59	0.59	0.59	107516
macro avg	0.44	0.43	0.41	107516
weighted avg	0.54	0.59	0.55	107516

```
-----
```

```
('Accuracy for the Multinomial Naive Bayes model:', 58.90751144015775, '%')  
-----
```

Challenges and issues:

We got an accuracy of 63.23%, which is the highest we have achieved so far. However, despite performing so well, Naïve Bayes model isn't that flexible when it comes to changing number of bins. When we tuned the number of bins from 3 to 5, the accuracy drops down massively to

48.46%. Now, comparing on a plain level field, Naïve Bayes for 5 bins has an accuracy of 48.46%, while, Random Forest Classifier for 5 bins had the highest accuracy of 62.18%, thus making it the better model. Apart from this issue, we didn't face any other challenge or issue while implementing Naïve Bayes model. It is mostly because most of the procedure mirrors that of the Random Forest Classifier (pre-processing) and only the modeling code changed, which was pretty straight forward.

```
-----  
Enter the number of bins: 5  
-----
```

```
('For', 5, 'bins, bin width=', 4755)  
-----
```

```
[1, 1, 1, 2, 1, ..., 3, 1, 2, 2, 4]  
Length: 20  
Categories (5, int64): [0 < 1 < 2 < 3 < 4]  
-----
```

```
-----  
The confusion matrix for the Gaussian Naive Bayes model is as follows:
```

```
[[ 4911 10208 1179 1367 142]  
 [ 2111 38352 2655 4063 3]  
 [ 969 11821 3123 4848 8]  
 [ 1259 5221 3298 5709 9]  
 [ 673 2397 1214 1970 6]]
```

```
-----  
The classification report for the Gaussian Naive Bayes model is as follows:
```

	precision	recall	f1-score	support
0	0.49	0.28	0.35	17807
1	0.56	0.81	0.67	47184
2	0.27	0.15	0.19	20769
3	0.32	0.37	0.34	15496
4	0.04	0.00	0.00	6260
micro avg	0.48	0.48	0.48	107516
macro avg	0.34	0.32	0.31	107516
weighted avg	0.43	0.48	0.44	107516

```
-----  
('Accuracy for the Gaussian Naive Bayes model:', 48.45883403400424, '%')  
-----
```

```
The confusion matrix for the Multinomial Naive Bayes model is as follows:
```

```
[[ 1474 12662 449 2833 389]  
 [ 605 38451 774 7158 196]  
 [ 665 10947 1272 7784 101]  
 [ 782 6066 226 8332 90]  
 [ 424 2716 110 2911 99]]
```

```
-----  
The classification report for the Multinomial Naive Bayes model is as follows:
```

	precision	recall	f1-score	support
0	0.37	0.08	0.14	17807
1	0.54	0.81	0.65	47184
2	0.45	0.06	0.11	20769
3	0.29	0.54	0.37	15496
4	0.11	0.02	0.03	6260
micro avg	0.46	0.46	0.46	107516
macro avg	0.35	0.30	0.26	107516
weighted avg	0.43	0.46	0.38	107516

```
-----  
('Accuracy for the Multinomial Naive Bayes model:', 46.15871126157967, '%')  
-----
```

Support Vector Machine Model:

Another algorithm we used here is Support Vector Machine. SVM algorithm maximizes the margin between the two separating hyperplanes by finding the maximum of the functional. In addition, SVM could solve both linearly separable and inseparable problems by implementing different kernels. Because we do not know how our data set is separated, we used both Linear kernel and RBF kernel to test which one is better.

Another problem we faced here is that running SVM usually takes lots of time. In order to accelerate the processing time, we use two algorithms here. First is Bagging Classifier, A Bagging classifier is an ensemble meta-estimator that fits base classifiers each on random subsets of the original dataset and then aggregates their individual predictions (either by voting or by averaging) to form a final prediction. By using Bagging Classifier, it could reduce the error and avoid model overfitting. Second method is One Vs Rest Classifier, this strategy consists in fitting one classifier per class. For each classifier, the class is fitted against all the other classes. In addition to its computational, one advantage of this approach is its interpretability. Since each class is represented by one and one classifier only, it is possible to gain knowledge about the class by inspecting its corresponding classifier. By implementing those two methods together, we could get a better result with high accuracy and low time cost. The data processing steps for SVM models are the same as Random Forest model, we encode the categorical variables into numbers and scale the numerical variables. The bins we set here are 5 bins, which equally separate purchase amount into 5 intervals. 20% of original data as test data set, and 80% as train data set.

Linear Kernel:

```
[[ 710 16385    22   166   524]
 [ 1973 42620    25   644  1922]
 [ 2164 16029    33  1040  1503]
 [ 2550  9833    49  1353  1711]
 [  859  4327    12   448   614]]
```

Confusion Matrix

```
In [186]: print(accuracy_score(y_test_bin,pre))
0.4216116671007106
```

Accuracy

The rows of confusion matrix are the instances which model predicts for one to five bins. The columns are the original instances in test data set. For example, on the first columns, 710 means that at bin one, there are 710 instances are correctly predicting as bin one, 1973 instances are

predicting as bin two, 2164 are predicting as bin three, 2550 are predicting as bin four, 859 are predict as bin five. Hence, we could find that bin two has the best prediction. Bin two has the worst prediction result. The overall accuracy for Linear kernel method is 42.16%. The time consumption of this model is 2023 seconds.

RBF Kernel:

```
[[ 3281 12547  428 1463  88]
 [ 1034 42162  903 2995  90]
 [  850 11297 1875 6619 128]
 [ 1003 2917 1087 10269 220]
 [  553 2034  205 3396  72]]
```

Confusion Matrix

```
In [189]: print(accuracy_score(y_test_bin,pre_1))
0.5362829718367499
```

Accuracy

In here, we could find that the RBF kernel gives us a much higher accuracy than linear kernel. This also indicate that our data set is not linearly separable. Because we separate our data randomly each time, so there are some differences in bin numbers, but we could still find that bin 2 has the highest accuracy result. However, the time cost of RBF kernel model is 6864 seconds, which is three times higher than Linear Kernel, this is because the RBF kernel is more complicated than Linear kernel.

Comparison and Summary of Models:

Linear regression: Adjusted R-squared: 0.3 ,16 out of 19 parameters are significant, 5.88 seconds on time consumption.

Random Forest Regressor: smallest mean squared error 2902, maximum number of leaf nodes 2000, 11.36 seconds on time consumption.

Random Forest Classifier: accuracy rate 62.18% with 5 bins, Number of tress 150, maximum leaf nodes 1500, entropy criteria, 90.13 seconds on time consumption.

Naïve Bayes: accuracy rate 63.23% with 3 bins 5.36 seconds on time consumption.

SVM: accuracy rate 53.6% with 5 bins. 6864 seconds on time consumption, RBF Kernel.

Apparently, Naïve Bayes model has the highest accuracy. However, this rate is achieved when the target bins is 3. If we increase the bins to 4, the accuracy drops severely down to 48.46%. because the when the number of bins increase, the number of borders is increasing. This could cause incorrectly classifying. In addition, if we only cut the purchase amount into 3 bins, the prediction will be too broad. In this case, we want to predict the purchase amount of a consumer as accurate as possible, so ideally, we want to use as many as bins as we could. Hence the Naïve bayes model is not the best choice. In summary, if we consider the time cost as a bias, the random forest classifier is our best model here, it has a 62% accuracy rate with 90 seconds of time cost. Although it has the second highest processing time, the 90 second time consumption is not beyond our tolerance. If we change the number of bins to 4 or 5, it will not have too much drop on accuracy. However, if we do not care about the time cost, the SVM model will be our first choice. The SVM model has accuracy of 54% with 5 bins. In addition, if we want to predict a certain number of purchase amount instead of a range of purchase amount, the random forest regressor will be the best model.

User interface and flexibility:

```
-----  
Welcome to the system.  
1). Press 1 for Linear regression model.  
2). Press 2 for Random Forest Regressor.  
3). Press 3 for Random Forest Classifier.  
4). Press 4 for Gaussian Naïve Bayes.  
5). Press 5 to exit.  
Your choice: 5  
Goodbye!  
(base) Vijets-MacBook-Air:IST707 vijetmuley$ █
```

Random Forest Regressor:

```
-----  
Maximum number of leaf nodes you wish to have: 2000  
-----
```

Random Forest Classifier:

```
-----  
Welcome to the system.  
1). Press 1 for Linear regression model.  
2). Press 2 for Random Forest Regressor.  
3). Press 3 for Random Forest Classifier.  
4). Press 4 for Gaussian Naive Bayes.  
5). Press 5 to exit.  
Your choice: 4  
-----
```

```
Enter the number of bins: 3  
-----
```

```
-----  
Maximum number of leaf nodes you wish to have: 1250  
-----
```

```
Welcome to the system.  
1). Press 1 for Linear regression model.  
2). Press 2 for Random Forest Regressor.  
3). Press 3 for Random Forest Classifier.  
4). Press 4 for Gaussian Naive Bayes.  
5). Press 5 to exit.  
Your choice: 3  
-----
```

```
Enter the number of bins: 5  
-----
```

In order to make our programming user friendly, we design a simple user interface which allow them select model. For specific model such random forest, they could select bins and leaf nodes they want to use. Finally, the program will return the time cost and accuracy rate of model the select.

Conclusion:

In this project, our goal is to find the best machine learning model which could give us a best predicting result. The biggest problem we have is how to format the data set to fit into our models. We spend lots time on data cleaning and re-formatting the data. Finally. We successfully implement 5 models to learn patterns from the data and help in forecasting sales amount. The Random Forest Classifier is the best model if we want to predict a range of purchase amount of a consumer. The Random Forest Regression can be used when we want to find a specific amount of purchase of a consumer. The user interface panel offers other people a easy and flexible way to do the analysis they want with our program.